

6. Neural audio compression

Audio coding · neural codecs, CNNs and autoencoders

Key question: can a neural network learn the analysis, quantization and reconstruction steps of an audio codec?

Learning goals

- Describe the encoder–quantizer–decoder structure of neural audio codecs.
- Explain why CNNs are useful for learning audio representations.
- Relate convolution, pooling, bottlenecks and autoencoders to compression.
- Explain residual vector quantization as the discrete coding stage.
- Compare reconstruction, perceptual and adversarial losses.
- Discuss current limits: compute, latency, generalization and evaluation.

Roadmap

1 · Neural codecs

SoundStream
EnCodec
Results and bitrates

2 · CNN building blocks


Convolution
Kernels, stride, padding
Pooling and hierarchy
Convolutional autoencoders

3 · Training and limits

Bottleneck and losses
RVQ and transformers
PyTorch implementation
Current challenges


Compression is still a rate–distortion problem, but many parts of the representation are learned from data.

SoundStream



Training combines reconstruction and adversarial losses; inference uses the encoder, RVQ codes and decoder.

EnCodec



Key idea: reduce a 24 kHz audio waveform to a low-rate discrete token stream, then reconstruct perceptually plausible audio.


Reported compression results

Model	Bandwidth	Entropy Coded	Compression	MUSHRA
Reference	-	-	1×	95.1±1.8
MP3	64 kbps	-	24×	82.7±3.2
Opus	6 kbps	-	256×	17.7±5.9
Opus	24 kbps	-	64×	82.9±3.7
ENCODER	6 kbps	4.2 kbps	256×	82.9±3.4

MUSHRA scores compare perceived quality: the same nominal bitrate can sound very different depending on the model.

Convolutional Neural Networks (CNNs)


- Originally developed for images but widely applied to audio.
- Learn feature extractors automatically, from low-level to high-level structure.
- For audio, filters may capture transients, timbre, envelopes or local spectral patterns.
- Key strength: automatic feature extraction (from low-level to high-level).



In neural audio coding, the “analysis filter bank” is no longer fixed: it is learned.

What is convolution?

- Mathematical operation between a filter/kernel and the input signal.
- Each position computes a weighted sum over a local neighborhood.
- Unlike classical filters, CNN kernels are learned from training data.



For waveform audio, 1D convolution is a natural way to learn time-local analysis filters.

Kernels and feature detection

Learned filters

Small matrices or 1D kernels
Examples: 3×3 , 5×5 , or short waveform kernels

Audio examples

Low-frequency energy
High-frequency detail
Attacks and transients
Timbre-related local structure

Training

The useful filters are learned from data during optimization.

The model decides which local patterns matter for reconstruction and perceptual quality.

Stride and padding in CNNs

Stride

How much the filter moves between positions.

Larger stride reduces the number of output frames and can contribute to compression.

Padding

Extra samples around the input, often zeros.

Padding controls the output size and avoids losing boundary information.

In learned codecs, strided convolutions often replace explicit downsampling stages.

Convolutional layer: extracting features

- Each filter produces one feature map.
- Each feature map highlights a specific local pattern.
- Stacked layers capture progressively more complex temporal and spectral structure.

Layer 1: local waveform patterns


Middle layers: timbre / envelopes

Deep layers: compact latent code

Compression emerges when the latent representation is smaller and quantized.

Non-linear activation functions


Commonly used: ReLU keeps positive values and sets negative values to zero.



Without non-linearities, a deep network would collapse into a mostly linear transform.

Pooling layer: downsampling the data

- Max-pooling takes the maximum value in a local window.
- It reduces dimensionality and computation.
- It can make the representation less sensitive to small time shifts.



For audio coding, downsampling must preserve the information needed for perceptual reconstruction.

How CNNs build hierarchical representations

Early layers

Edges, attacks,
transients, local
oscillations

Deeper layers

Harmonics, timbre,
envelopes, longer
temporal dependencies

End-to-end training

The model is trained
from waveform input to
compressed
representation and
output.

This is the learned counterpart of hand-designed filter banks and transforms.

CNNs vs. traditional audio filters

Traditional audio coding

- Predefined filters and transforms
- Examples: Fourier, MDCT, filter banks
- Designed from signal-processing theory


Neural audio coding

- Learned filters and latent spaces
- Optimized for data, loss function and bitrate
- Can adapt to speech, music or mixed audio

The gain depends on whether the learned representation generalizes beyond the training distribution.

Convolutional Autoencoder (CAE)

- Unsupervised or self-supervised learning objective.
- Encoder compresses data; decoder reconstructs data.
- In audio, the compact latent representation becomes the object to quantize and code.



The bottleneck is the architectural equivalent of the bit-rate constraint.

How the encoder works in audio compression

- Series of convolutional layers compresses the input audio.
- CNN layers learn features such as harmonics, transients and envelopes.
- The final encoder output is a bottleneck: a compact version of the input.

input waveform → learned analysis → latent representation

How the decoder works in audio compression

- Uses transposed convolutions or related upsampling layers.
- Reconstructs the original signal as accurately and perceptually as possible.
- Success depends on whether the encoder retained the essential audio features.

latent code → learned synthesis → reconstructed waveform


Understanding the bottleneck

Smaller bottleneck

- Higher compression
- Lower bitrate
- Higher risk of losing audible detail

Larger bottleneck

- Better reconstruction
- More latent information
- Higher bitrate and storage cost



The codec design problem is finding the best quality–rate balance for the use case.

Training the autoencoder

- Signal loss: often Mean Squared Error (MSE) or waveform/spectral distance.
- Perceptual loss: encourages differences that are less audible.
- Optimization uses backpropagation and gradient descent.

The loss function defines what “good reconstruction” means.

Gradient descent and learning rate

Gradient descent

- Iteratively moves parameters in the direction that reduces the loss.
- This is the core optimization loop for training the encoder, quantizer support modules and decoder.

Learning rate

- Controls the step size of parameter updates.
- Too large can destabilize training; too small can make training very slow.

Neural codec quality depends on both architecture and training dynamics.

The discriminator and perceptual loss

- Traditional losses like MSE penalize all differences, even inaudible ones.
- A discriminator learns to distinguish real from reconstructed audio.
- The generator is pushed toward statistically and perceptually plausible outputs.

Risk: better perceived realism may not always mean faithful reconstruction.

Residual Vector Quantization (RVQ)

1. Map the input vector to the nearest code vector in a learned codebook.
2. Compute the residual between the input and the selected code vector.
3. Quantize the residual with another codebook.
4. Repeat across several layers to refine the approximation.

RVQ is the key quantization mechanism used by SoundStream and EnCodec.

Transformer layer: context through attention

- Self-attention weighs the relevance of other positions when processing one frame.
- This can model token statistics produced by the RVQ codebooks.
- Better token modeling can help entropy coding assign shorter bitstrings to common patterns.

Two sub-layers

Multi-head self-attention
Feed-forward network

The transformer is not the waveform codec itself; it helps exploit context in the token sequence.

PyTorch implementation

- Open-source machine learning framework with a Pythonic interface.
- GPU acceleration is crucial for training large audio models.
- TorchAudio provides audio-specific operations such as spectrograms, mel-spectrograms and resampling.

A typical class demo can train a small convolutional autoencoder before moving to RVQ-based systems.

Challenges and current limitations

Compute and latency

Large architectures require compute power.

Low-latency real-time communication remains difficult.

Generalization

Models may struggle across speech, music and environmental audio.

Training needs large and diverse datasets.

Evaluation

Subjective tests such as MUSHRA are expensive and difficult to scale.

The best neural codec is not just the best model: it must be efficient, robust and audibly reliable.

Key takeaways

- Neural audio codecs learn the analysis and synthesis transforms instead of fixing them by design.
- The bottleneck and RVQ layers determine the discrete, low-rate representation.
- Perceptual and adversarial losses target audible quality, not only numerical error.
- CNNs provide local feature extraction; transformers can model longer token context.
- Compute, latency, generalization and listening evaluation are still central challenges.

References and notebook examples

- Neil Zeghidour, Alejandro Luebs, Ahmed Omran, Jan Skoglund, Marco Tagliasacchi (2021). SoundStream: An End-to-End Neural Audio Codec.
- Alexandre Défossez, Jade Copet, Gabriel Synnaeve, Yossi Adi (2022). High Fidelity Neural Audio Compression.
- PyTorch: pytorch.org
- Course repository: github.com/xserra/audio-coding-materials/
 - [examples/convolutional-autoencoder-simple.ipynb](#)
 - [examples/convolutional-autoencoder-complex.ipynb](#)